

# TASK TWO: PREDICTIVE MODELING WITH LINEAR REGRESSION

Implement a simple linear regression model using a dataset with continuous target variables. Split the data into training and testing sets, train the model on the training data, evaluate its performance using metrics like mean squared error or R-squared, and make predictions on the test set. Visualize the regression line and actual vs. predicted values to assess the model's accuracy.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [2]: ## Load the dataset
df = pd.read_csv('petrol_consumption.csv')
```

```
In [3]: # Show the first few rows and the structure of the dataset
print(df.head())
```

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%) \
0	9.0	3571	1976	0.525
1	9.0	4092	1250	0.572
2	9.0	3865	1586	0.580
3	7.5	4870	2351	0.529
4	8.0	4399	431	0.544

	Petrol_Consumption
0	541
1	524
2	561
3	414
4	410

```
In [4]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48 entries, 0 to 47
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Petrol_tax                            48 non-null     float64
1   Average_income                        48 non-null     int64
2   Paved_Highways                        48 non-null     int64
3   Population_Driver_licence(%)          48 non-null     float64
4   Petrol_Consumption                    48 non-null     int64
dtypes: float64(2), int64(3)
memory usage: 2.0 KB
None
```

## Define Features and Target Variable

```
In [5]: # Define feature matrix (X) and target variable (y)
X = df[['Petrol_tax', 'Average_income', 'Paved_Highways', 'Population_Driver_licence(%)']]
y = df['Petrol_Consumption']
```

```
In [6]: X
```

Out[6]:

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%)
0	9.00	3571	1976	0.525
1	9.00	4092	1250	0.572
2	9.00	3865	1586	0.580
3	7.50	4870	2351	0.529
4	8.00	4399	431	0.544
5	10.00	5342	1333	0.571
6	8.00	5319	11868	0.451
7	8.00	5126	2138	0.553
8	8.00	4447	8577	0.529
9	7.00	4512	8507	0.552
10	8.00	4391	5939	0.530
11	7.50	5126	14186	0.525
12	7.00	4817	6930	0.574
13	7.00	4207	6580	0.545
14	7.00	4332	8159	0.608
15	7.00	4318	10340	0.586
16	7.00	4206	8508	0.572
17	7.00	3718	4725	0.540
18	7.00	4716	5915	0.724
19	8.50	4341	6010	0.677
20	7.00	4593	7834	0.663
21	8.00	4983	602	0.602
22	9.00	4897	2449	0.511
23	9.00	4258	4686	0.517
24	8.50	4574	2619	0.551
25	9.00	3721	4746	0.544
26	8.00	3448	5399	0.548
27	7.50	3846	9061	0.579
28	8.00	4188	5975	0.563
29	9.00	3601	4650	0.493
30	7.00	3640	6905	0.518
31	7.00	3333	6594	0.513
32	8.00	3063	6524	0.578
33	7.50	3357	4121	0.547
34	8.00	3528	3495	0.487
35	6.58	3802	7834	0.629
36	5.00	4045	17782	0.566
37	7.00	3897	6385	0.586
38	8.50	3635	3274	0.663
39	7.00	4345	3905	0.672
40	7.00	4449	4639	0.626
41	7.00	3656	3985	0.563
42	7.00	4300	3635	0.603
43	7.00	3745	2611	0.508
44	6.00	5215	2302	0.672
45	9.00	4476	3942	0.571
46	7.00	4296	4083	0.623
47	7.00	5002	9794	0.593

In [7]:

y

```
Out[7]: 0      541
1      524
2      561
3      414
4      410
5      457
6      344
7      467
8      464
9      498
10     580
11     471
12     525
13     508
14     566
15     635
16     603
17     714
18     865
19     640
20     649
21     540
22     464
23     547
24     460
25     566
26     577
27     631
28     574
29     534
30     571
31     554
32     577
33     628
34     487
35     644
36     640
37     704
38     648
39     968
40     587
41     699
42     632
43     591
44     782
45     510
46     610
47     524
Name: Petrol_Consumption, dtype: int64
```

## Split the Data into Training and Testing Sets

```
In [8]: # Split the dataset into 80% training and 20% testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [9]: X_train
```

Out[9]:	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%)
8	8.00	4447	8577	0.529
3	7.50	4870	2351	0.529
6	8.00	5319	11868	0.451
39	7.00	4345	3905	0.672
33	7.50	3357	4121	0.547
13	7.00	4207	6580	0.545
17	7.00	3718	4725	0.540
45	9.00	4476	3942	0.571
15	7.00	4318	10340	0.586
9	7.00	4512	8507	0.552
16	7.00	4206	8508	0.572
29	9.00	3601	4650	0.493
32	8.00	3063	6524	0.578
46	7.00	4296	4083	0.623
0	9.00	3571	1976	0.525
31	7.00	3333	6594	0.513
30	7.00	3640	6905	0.518
5	10.00	5342	1333	0.571
11	7.50	5126	14186	0.525
34	8.00	3528	3495	0.487
1	9.00	4092	1250	0.572
44	6.00	5215	2302	0.672
21	8.00	4983	602	0.602
2	9.00	3865	1586	0.580
36	5.00	4045	17782	0.566
35	6.58	3802	7834	0.629
23	9.00	4258	4686	0.517
41	7.00	3656	3985	0.563
10	8.00	4391	5939	0.530
22	9.00	4897	2449	0.511
18	7.00	4716	5915	0.724
47	7.00	5002	9794	0.593
20	7.00	4593	7834	0.663
7	8.00	5126	2138	0.553
42	7.00	4300	3635	0.603
14	7.00	4332	8159	0.608
28	8.00	4188	5975	0.563
38	8.50	3635	3274	0.663

```
In [10]: y_train
```

```
Out[10]: 8      464
          3      414
          6      344
          39     968
          33     628
          13     508
          17     714
          45     510
          15     635
          9      498
          16     603
          29     534
          32     577
          46     610
          0      541
          31     554
          30     571
          5      457
          11     471
          34     487
          1      524
          44     782
          21     540
          2      561
          36     640
          35     644
          23     547
          41     699
          10     580
          22     464
          18     865
          47     524
          20     649
          7      467
          42     632
          14     566
          28     574
          38     648
Name: Petrol_Consumption, dtype: int64
```

```
In [11]: X_test
```

```
Out[11]:
```

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%)
27	7.5	3846	9061	0.579
40	7.0	4449	4639	0.626
26	8.0	3448	5399	0.548
43	7.0	3745	2611	0.508
24	8.5	4574	2619	0.551
37	7.0	3897	6385	0.586
12	7.0	4817	6930	0.574
19	8.5	4341	6010	0.677
4	8.0	4399	431	0.544
25	9.0	3721	4746	0.544

```
In [12]: y_test
```

```
Out[12]: 27      631
          40      587
          26      577
          43      591
          24      460
          37      704
          12      525
          19      640
          4      410
          25      566
Name: Petrol_Consumption, dtype: int64
```

## Train the Linear Regression Model

```
In [14]: # Initialize and train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[14]: ▼ LinearRegression
LinearRegression()
```

# Make Predictions on the Test Set

```
In [15]: # Make predictions using the test set
y_pred = model.predict(X_test)
```

```
In [16]: y_pred
```

```
Out[16]: array([606.69266519, 673.77944169, 584.99149034, 563.53691024,
        519.05867235, 643.46100256, 572.89761422, 687.07703573,
        547.6093662 , 530.03762971])
```

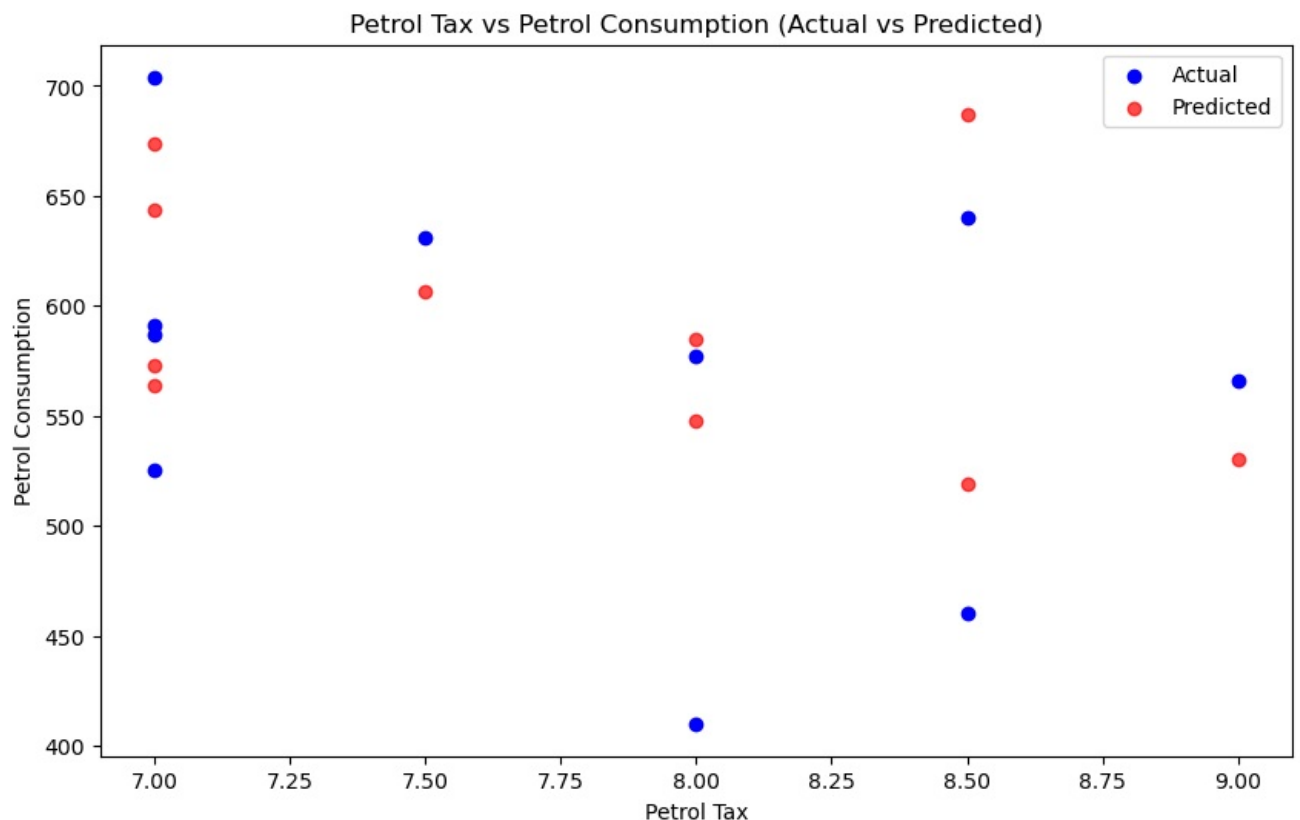
## Evaluate the Model (MSE and R-squared)

```
In [17]: # Calculate Mean Squared Error and R-squared score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')
```

```
Mean Squared Error: 4083.255871745411
R-squared: 0.3913664001428835
```

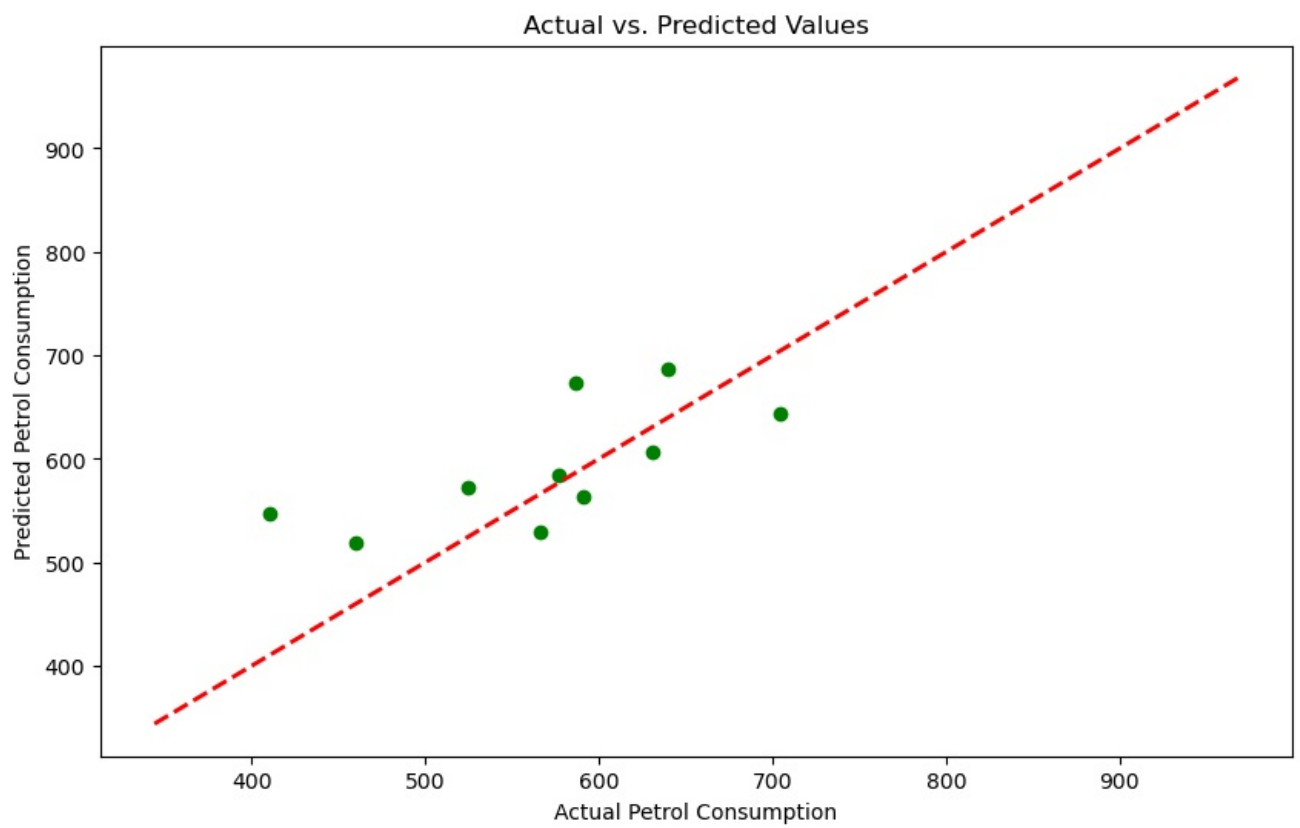
## Visualize the Regression Line (for One Feature)

```
In [18]: # Visualize the regression line for the 'Petrol_tax' feature
plt.figure(figsize=(10, 6))
plt.scatter(X_test['Petrol_tax'], y_test, color='blue', label='Actual')
plt.scatter(X_test['Petrol_tax'], y_pred, color='red', label='Predicted', alpha=0.7)
plt.title('Petrol Tax vs Petrol Consumption (Actual vs Predicted)')
plt.xlabel('Petrol Tax')
plt.ylabel('Petrol Consumption')
plt.legend()
plt.show()
```



## Visualize Actual vs Predicted Values

```
In [19]: # Visualize actual vs predicted petrol consumption
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='green')
plt.plot([y.min(), y.max()], [y.min(), y.max()], color='red', linestyle='--', linewidth=2)
plt.title('Actual vs. Predicted Values')
plt.xlabel('Actual Petrol Consumption')
plt.ylabel('Predicted Petrol Consumption')
plt.show()
```



## Coefficient Interpretation

```
In [20]: #Coefficients interpretation
coefficients = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
print(coefficients)
```

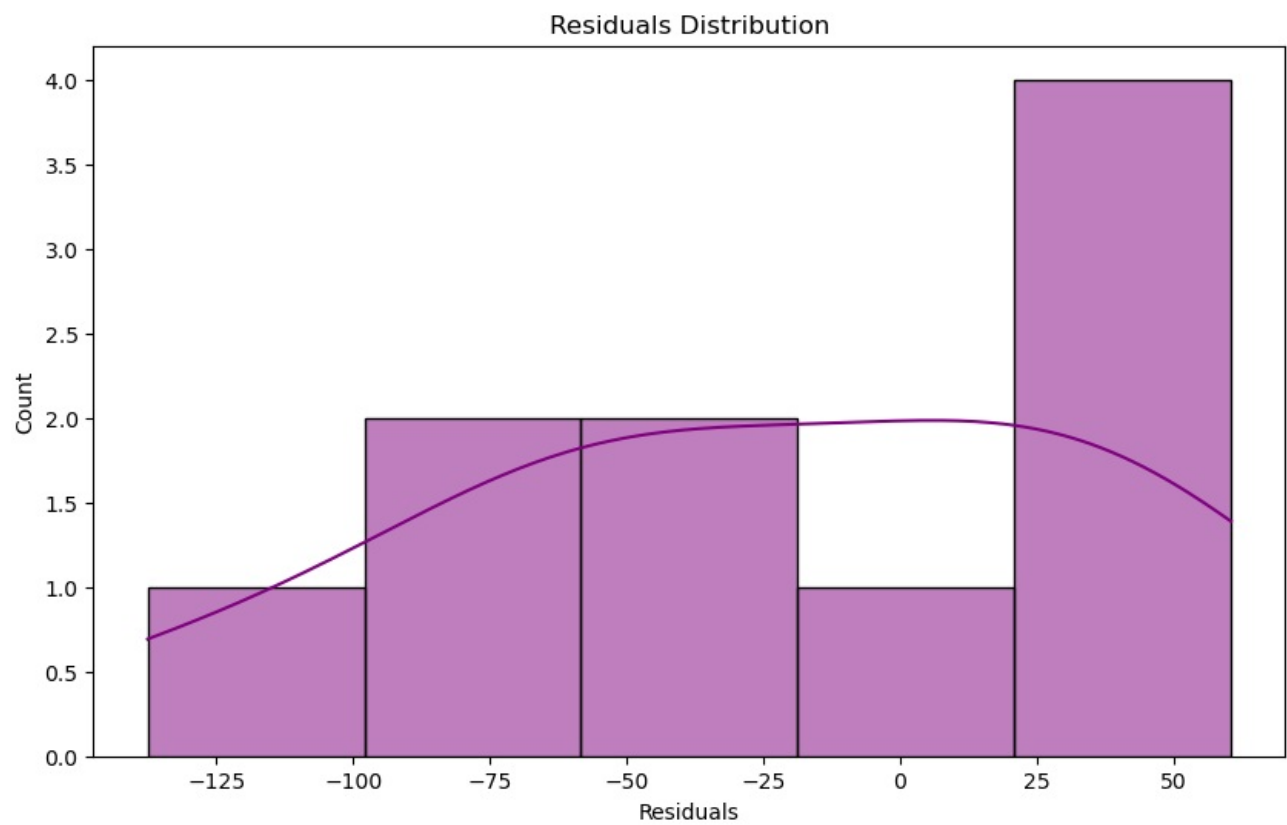
	Coefficient
Petrol_tax	-36.993746
Average_income	-0.056536
Paved_Highways	-0.004382
Population_Driver_licence(%)	1346.869298

## Residual Analysis

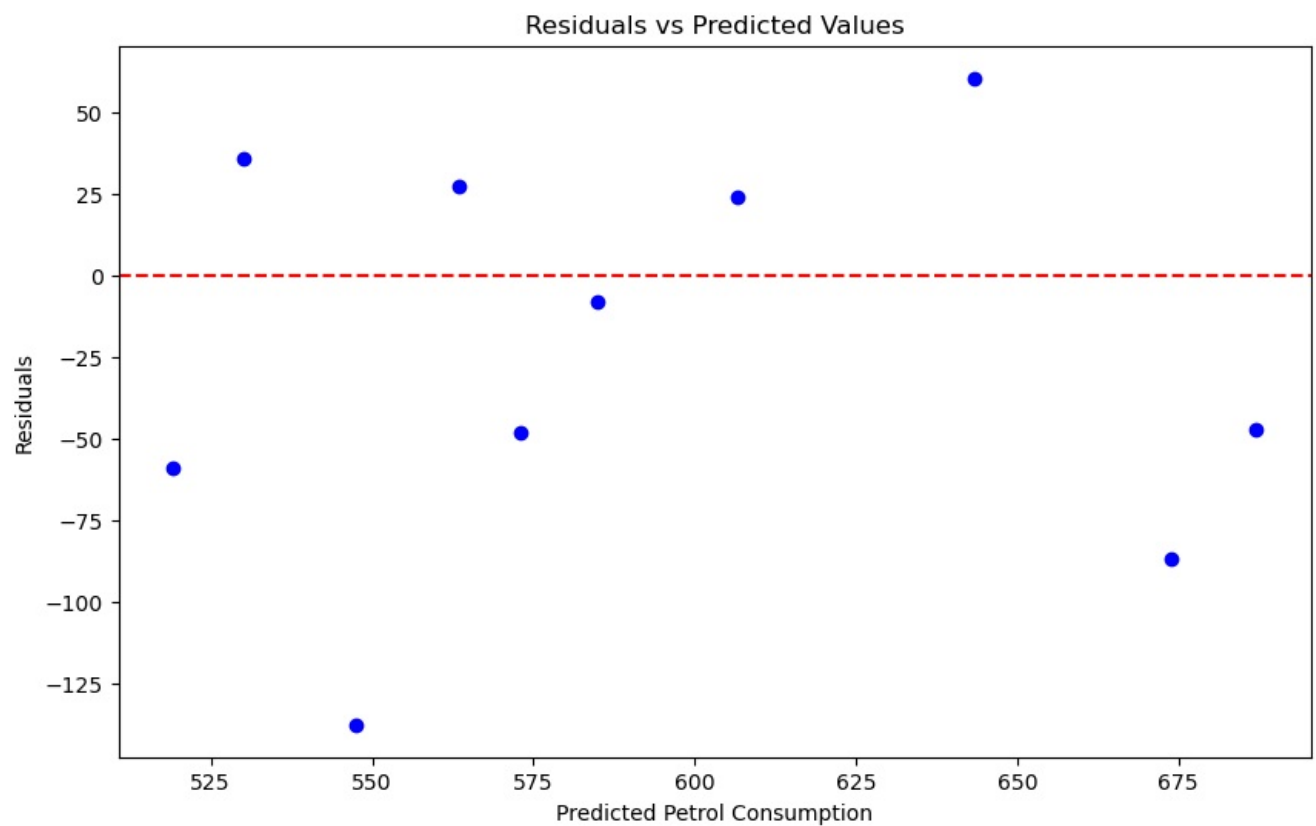
Residuals are the differences between actual and predicted values. Analyzing residuals helps check model assumptions.

```
In [21]: residuals = y_test - y_pred

# Plot the residuals
plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True, color='purple')
plt.title('Residuals Distribution')
plt.xlabel('Residuals')
plt.show()
```



```
In [22]: # Scatter plot of residuals vs. predicted values
plt.figure(figsize=(10, 6))
plt.scatter(y_pred, residuals, color='blue')
plt.axhline(y=0, color='red', linestyle='--')
plt.title('Residuals vs Predicted Values')
plt.xlabel('Predicted Petrol Consumption')
plt.ylabel('Residuals')
plt.show()
```



## Cross-Validation for Model Performance

Evaluate the model's robustness with cross-validation.

```
In [23]: from sklearn.model_selection import cross_val_score
```



```
# Cross-validation with 5 folds
cv_scores = cross_val_score(model, X, y, cv=5, scoring='neg_mean_squared_error')
mean_cv_score = np.mean(np.abs(cv_scores))

print(f'Mean Cross-Validation MSE: {mean_cv_score}')
```

Mean Cross-Validation MSE: 6012.505029227636

## Outlier Detection

Outliers can significantly affect your model's performance. Here's a way to detect and visualize outliers using Z-scores

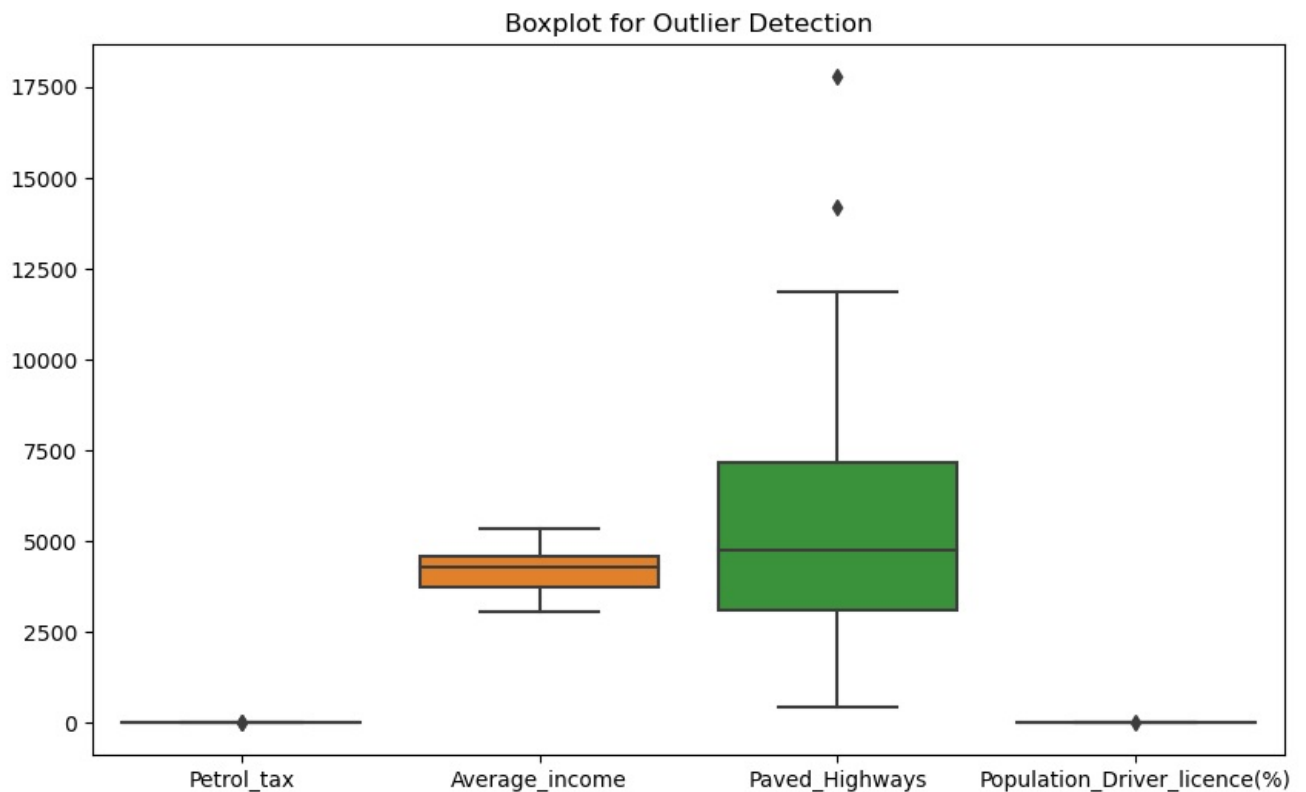
```
In [24]: from scipy import stats

# Calculate Z-scores for detecting outliers
z_scores = np.abs(stats.zscore(X))

# Filter out rows with Z-scores greater than 3
outliers = np.where(z_scores > 3)
print(f'Outliers found at indices: {outliers}')
```

```
# Visualize outliers in the dataset
plt.figure(figsize=(10, 6))
sns.boxplot(data=X)
plt.title('Boxplot for Outlier Detection')
plt.show()
```

Outliers found at indices: (array([36], dtype=int64), array([2], dtype=int64))



## Feature Scaling

Standardizing or normalizing your data ensures that features are on the same scale, which can improve model performance, especially for algorithms that rely on distance metrics.

```
In [26]: from sklearn.preprocessing import StandardScaler

# Apply standard scaling to the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the scaled data into training and testing sets
X_train_scaled, X_test_scaled, y_train_scaled, y_test_scaled = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Train a new linear regression model on the scaled data
model_scaled = LinearRegression()
model_scaled.fit(X_train_scaled, y_train_scaled)

# Evaluate the scaled model
y_pred_scaled = model_scaled.predict(X_test_scaled)
mse_scaled = mean_squared_error(y_test_scaled, y_pred_scaled)
```

```
r2_scaled = r2_score(y_test_scaled, y_pred_scaled)
print(f'Scaled Model MSE: {mse_scaled}')
print(f'Scaled Model R-squared: {r2_scaled}')
```

```
Scaled Model MSE: 4083.255871745399
Scaled Model R-squared: 0.39136640014288526
```

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js